

Persistent Data Structures

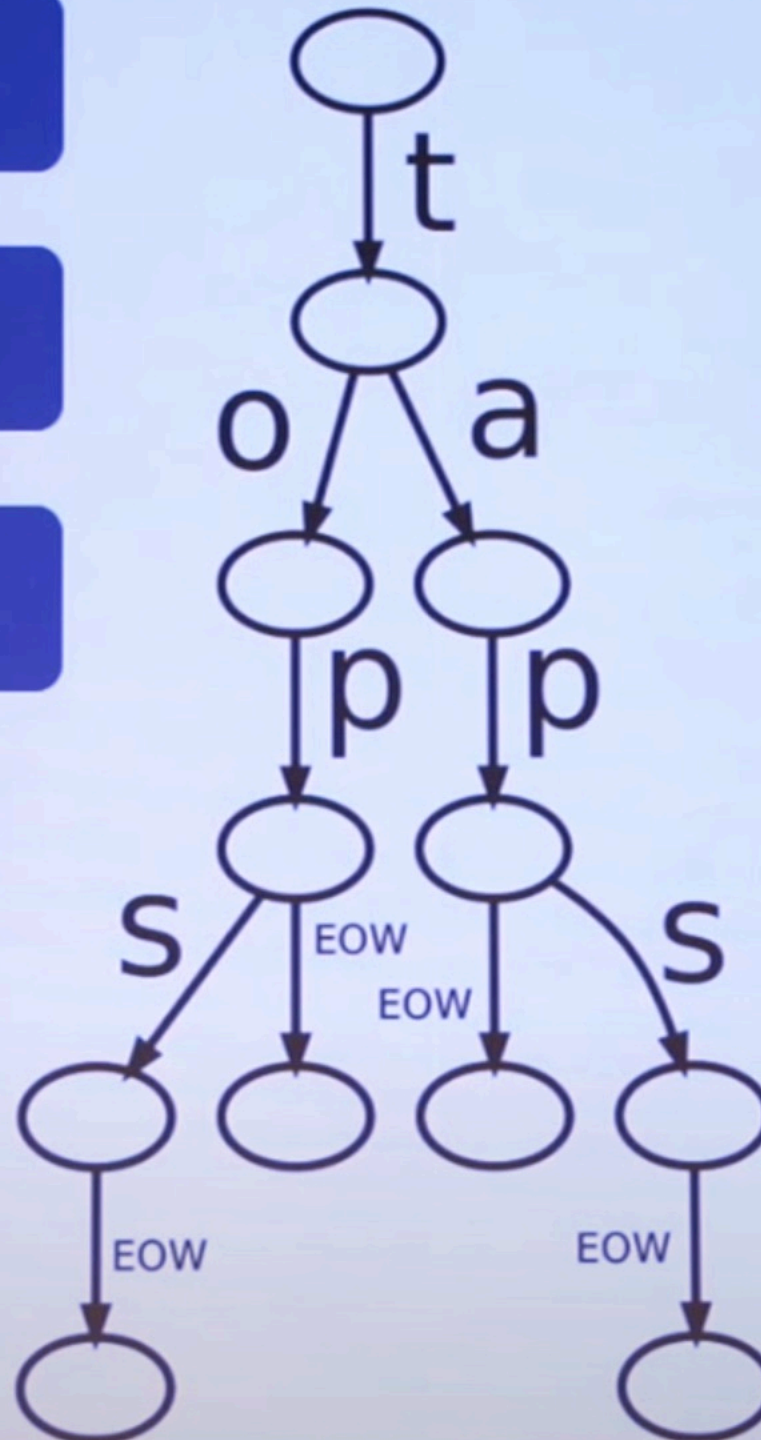
- Composite values - immutable
- 'Change' is merely a function, takes one value and returns another, 'changed' value
- Collection maintains its performance guarantees
 - Therefore new versions are not full copies
- Old version of the collection is still available after 'changes', with same performance
- Example - hash map/set and vector based upon array mapped hash tries (Bagwell)



LEVEL 0

LEVEL 1

LEVEL 2



00000 01011 01001 11001 10111 00010 11100

0

11

9

25

ArrayNode
shift = 0

28

ArrayNode
shift = 5

2

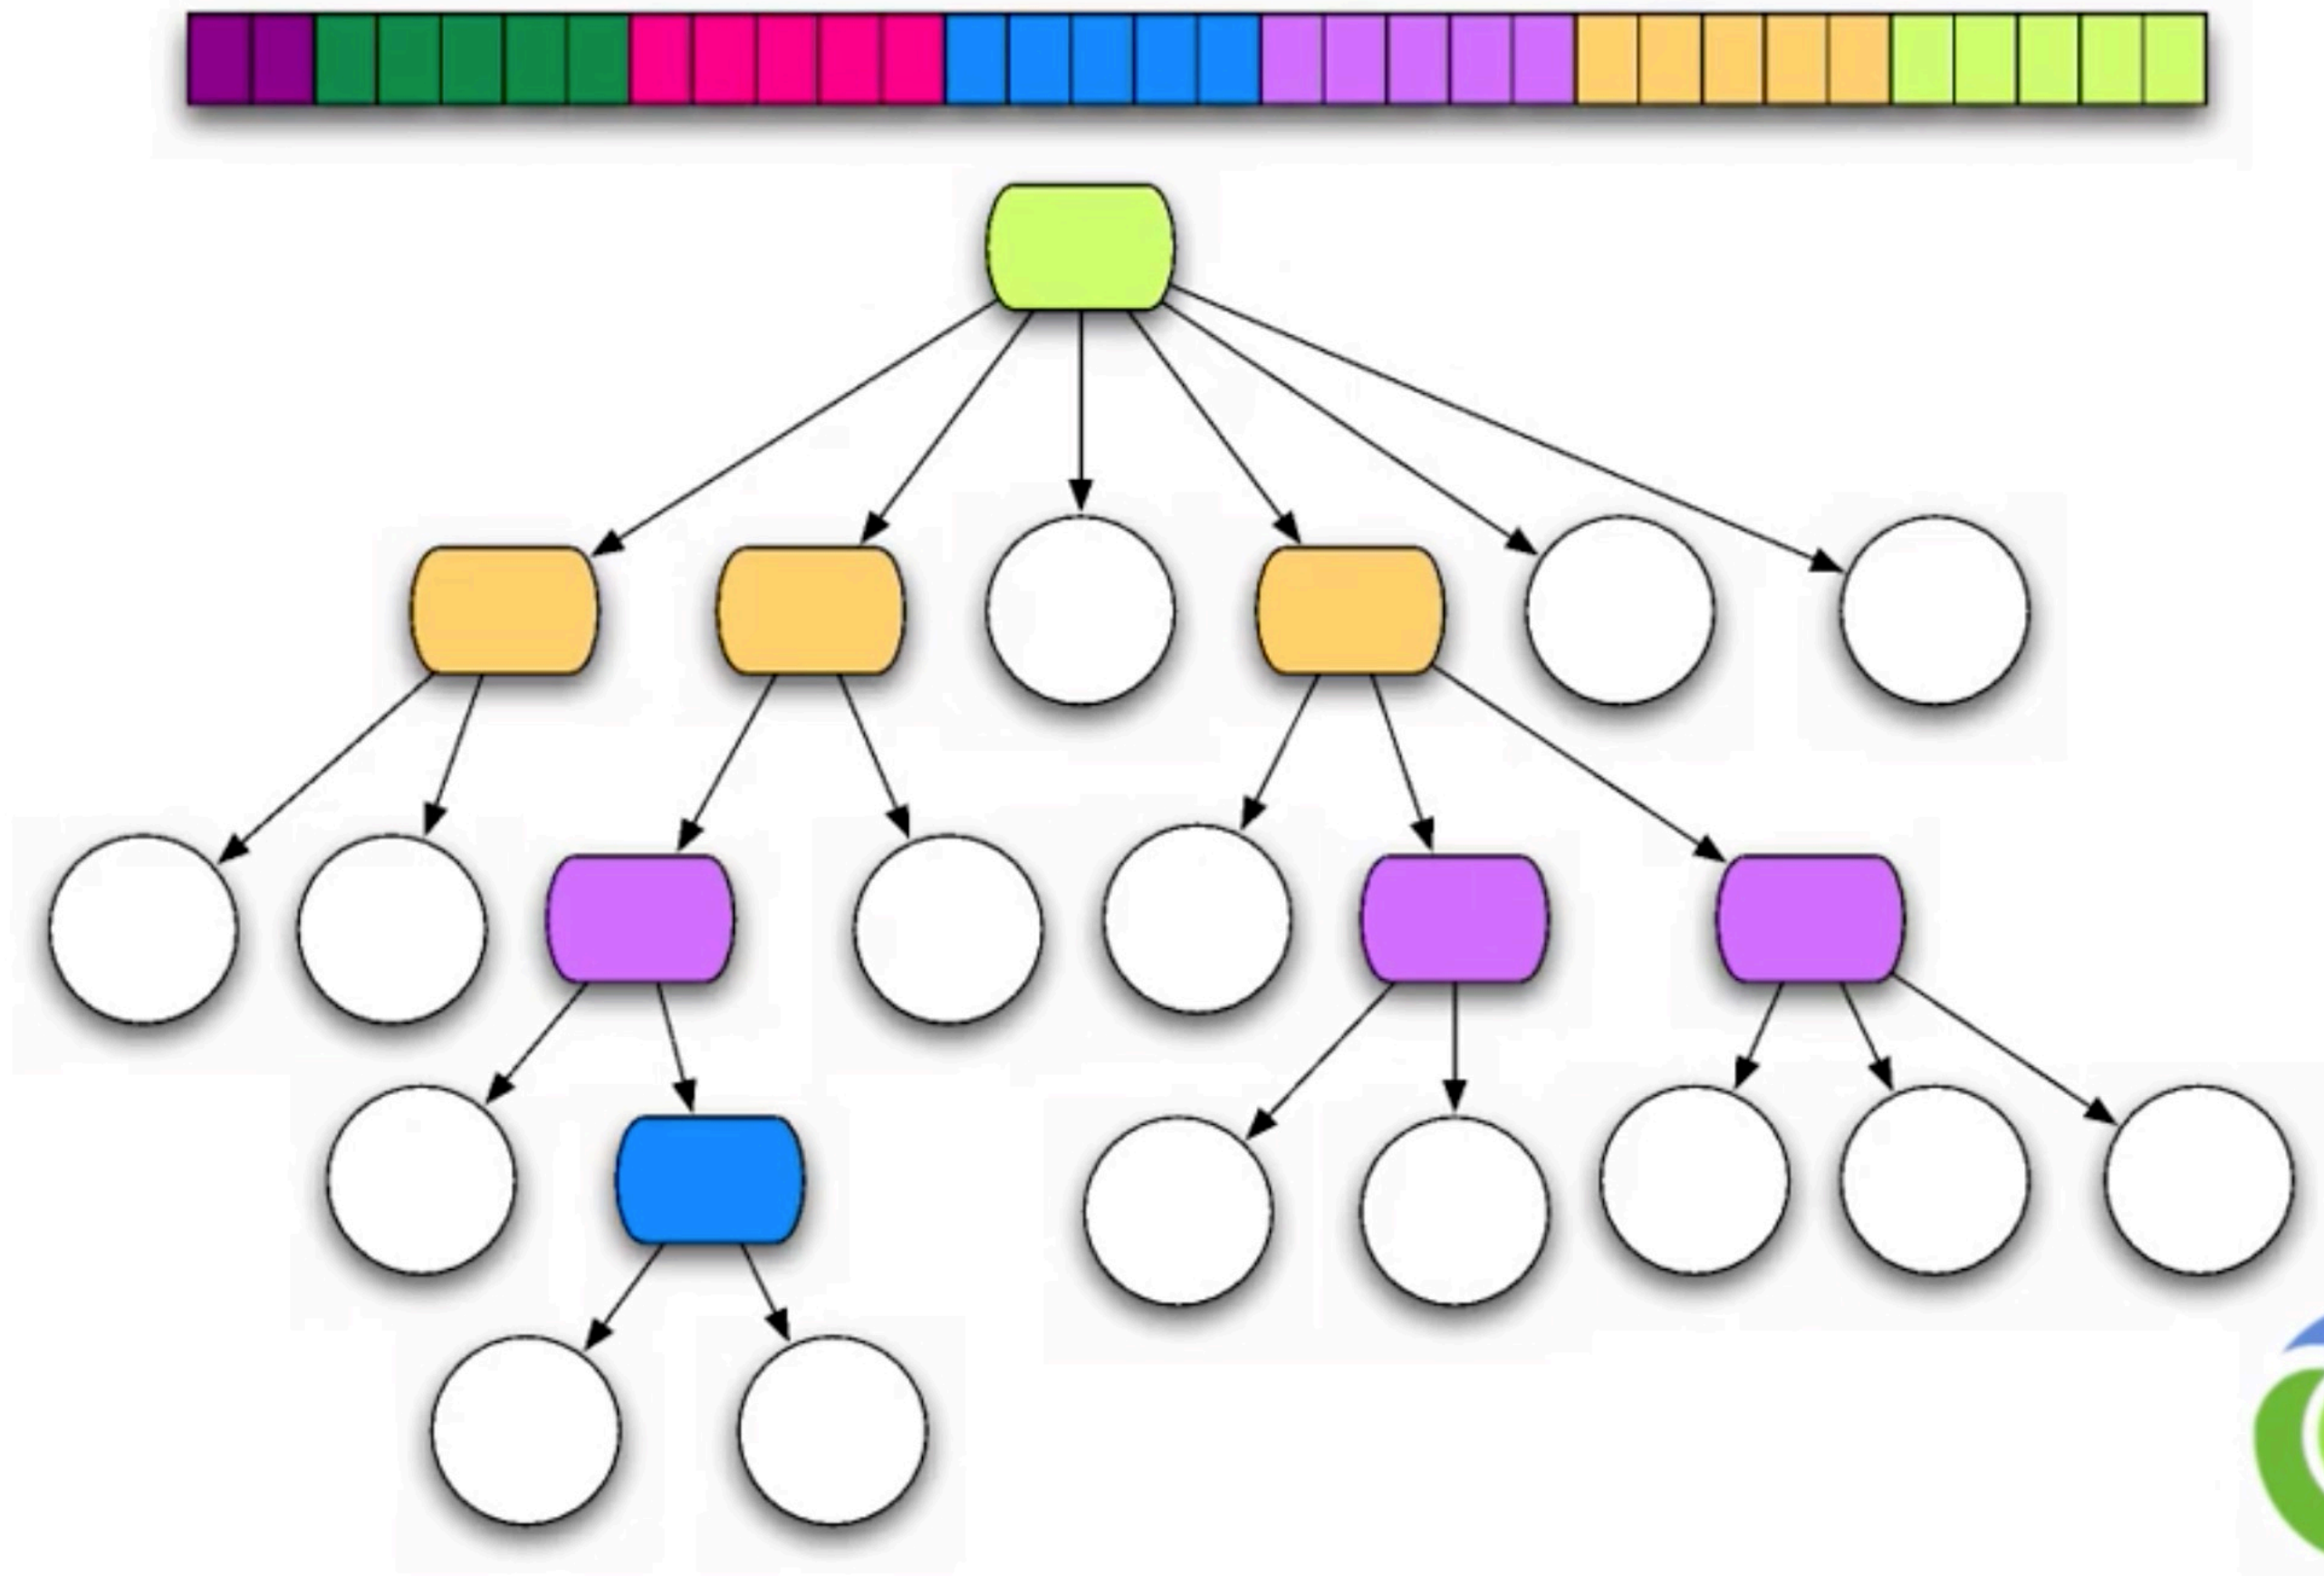
ArrayNode
shift = 10

23

BitmapIndexedNode
shift = 15

... and then follow the AMT down

Bit-partitioned hash tries



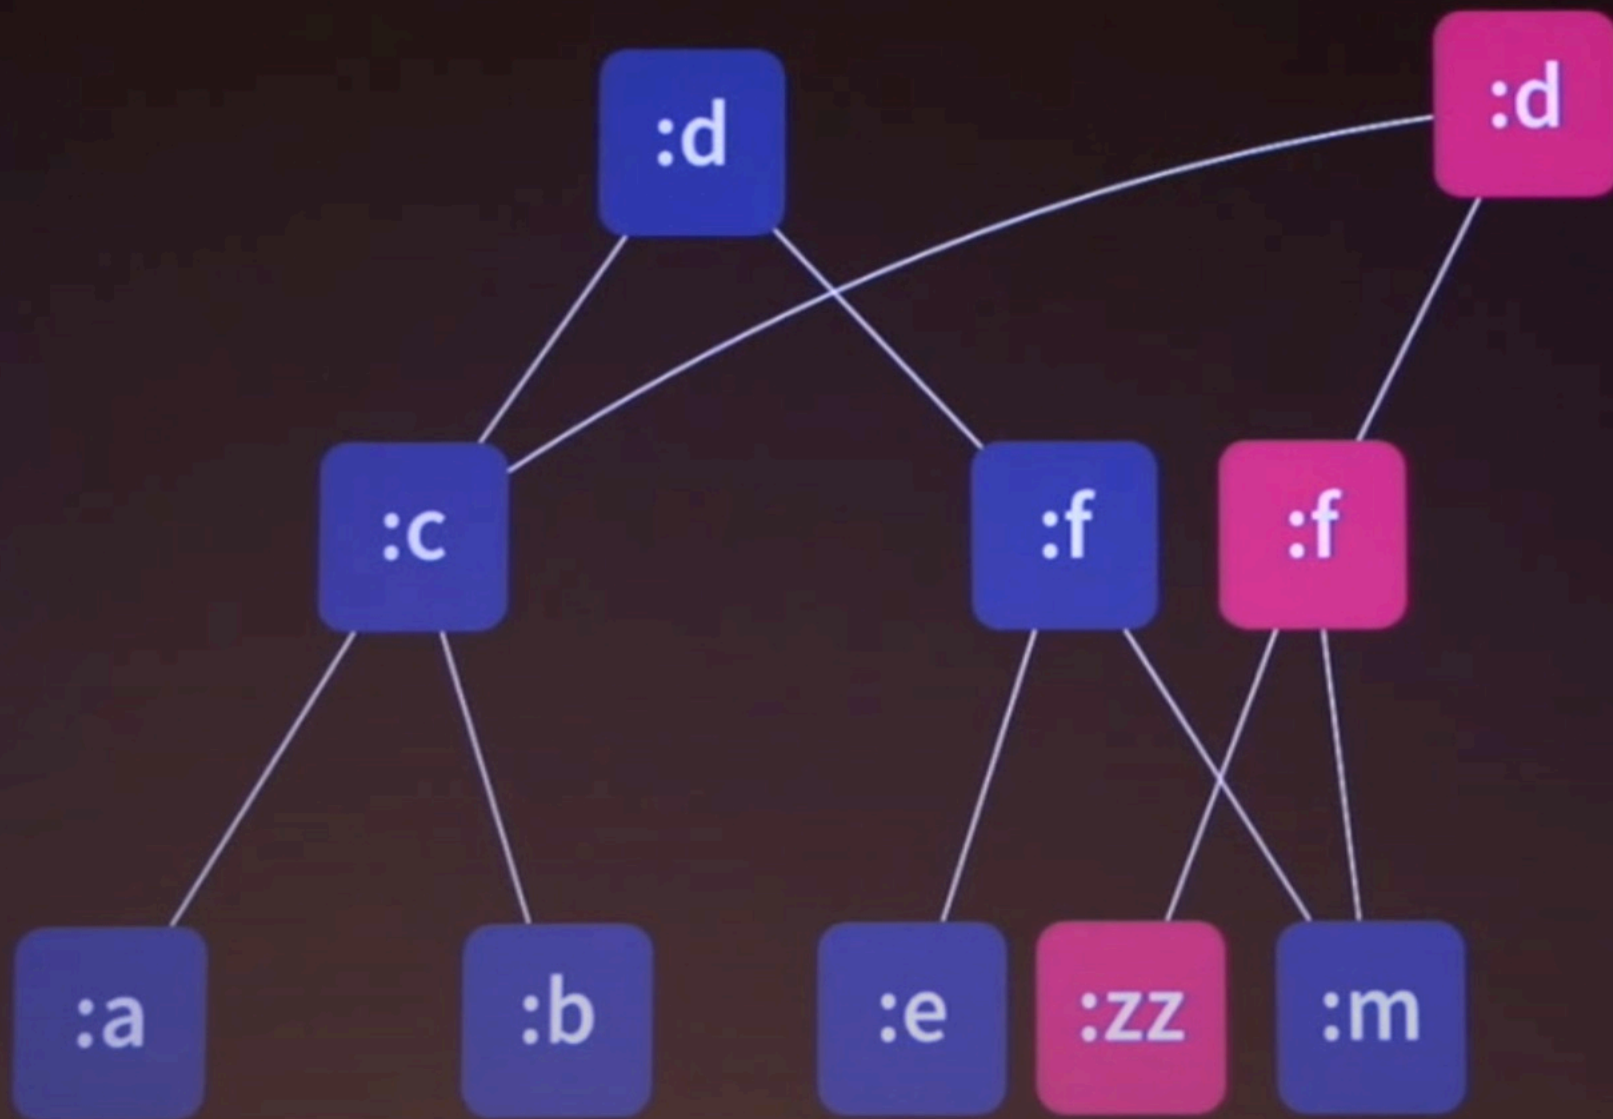
Structural Sharing

- Key to efficient 'copies' and therefore persistence
- Everything is immutable so no chance of interference
- Thread safe
- Iteration safe




STRUCTURAL SHARING

(assoc v 4 :zz)



Path Copying



REGULAR HASH TABLE?


NEED ROOT RESIZING

NOT AMENABLE TO
STRUCTURAL SHARING

THE TAIL OPTIMIZATION

PersistentVector

count shift root **tail**



Slides are taken from, the seminars:

1. What Lies Beneath - A Deep Dive Into Clojure's Data Structures - Mohit Thatte
2. Persistent Data Structures and Managed References - Rich Hickey